

Automated Pneumonia Detection from Chest Radiographs: Classical Machine Learning Architectures, Literature Synthesis, and Commercial Market Dynamics

Ahmed Ahmed Mokhtar Alhussien Ayman Hanafy
Ahmed Mohamed Abdelsalam

1 Introduction and Clinical Pathophysiology

Pneumonia remains one of the most lethal infectious diseases globally, responsible for an estimated 2.5 to 4 million deaths annually [1]. Its burden is borne most heavily by vulnerable populations: children under the age of five account for approximately 920,000 of those fatalities, while the elderly and immunocompromised face disproportionate mortality risk. In the United States alone, the disease generates over 500,000 emergency department visits and 50,000 deaths each year.

1.1 Pathophysiology and Radiographic Presentation

Pneumonia is characterised by an acute infection of the pulmonary parenchyma, caused by bacterial, viral, or fungal pathogens. The resulting inflammatory response engorges alveolar sacs with fluid and cellular debris, progressively compromising gas exchange and inducing hypoxia. Its clinical manifestation on chest radiographs (CXRs) varies markedly with the causative pathogen, as summarised in Table 1.

Pathogen Type	Radiographic Manifestation
Bacterial	Focal lobar consolidation; dense, contiguous radiopacity.
Viral (e.g., SARS-CoV-2)	Diffuse, bilateral interstitial pattern; ground-glass attenuation.

Table 1: Characteristic radiographic presentations by pneumonia subtype.

Manual radiographic interpretation is subject to well-documented systemic limitations. Inter-observer variability introduces meaningful diagnostic uncertainty, while the considerable visual overlap between pneumonia and other cardiopulmonary pathologies — including pleural effusions, pulmonary oedema, and lung carcinoma — further complicates accurate classification [2].

1.2 The Case for Classical Machine Learning

Convolutional Neural Networks (CNNs) have established state-of-the-art benchmarks in automated CXR diagnosis [2]. Nevertheless, classical machine learning frameworks retain indispensable advantages in clinical deployment contexts. Their reliance on handcrafted, mathematically derived features ensures algorithmic transparency, a prerequisite for clinical trust and regulatory acceptance. Furthermore, their computational lightness makes them robust on smaller

datasets and deployable in low-resource environments without specialised hardware acceleration — a critical consideration in regions where pneumonia mortality is highest.

2 Proposed System Architecture and Methodology

The proposed system implements a multi-stage classical machine learning pipeline for binary pneumonia classification, deliberately foregoing deep learning to preserve interpretability and deployment viability across resource-constrained settings.

2.1 Dataset and Statistical Profiling

The architecture is built upon the publicly available Chest X-Ray Images (Pneumonia) repository [3], comprising over 5,800 anterior-posterior paediatric radiographs. The dataset exhibits a severe class imbalance, with pneumonia instances constituting approximately 74% of the training corpus. The original 16-image validation partition was statistically negligible and was therefore merged into the training array. A new, stratified 80–20 split was generated, yielding approximately 1,043 validation samples while preserving the baseline class distribution.

2.2 Preprocessing and Standardisation Pipeline

All images undergo a three-step normalisation procedure prior to feature extraction. First, they are geometrically resampled to a uniform 64×64 pixel grid via bilinear interpolation, balancing structural detail against computational efficiency. Multi-channel (RGB/BGR) inputs are then converted to single-channel grayscale. Finally, pixel intensities are cast to `float32` and scaled to the unit interval using min-max normalisation:

$$x_{\text{norm}} = \frac{x - x_{\min}}{x_{\max} - x_{\min}} \quad (1)$$

2.3 Handcrafted Feature Engineering

To circumvent the overfitting that results from raw pixel flattening, the system computes an optimised vector of 89 distinct handcrafted features per radiograph across four complementary categories.

First-order statistical features characterise the global pixel distribution via the mean, variance, skewness (a useful indicator of unilateral opacities), kurtosis, and the 25th, 50th, and 75th percentiles.

Second-order textural features are extracted through the Gray-Level Co-occurrence Matrix (GLCM) [4], capturing micro-textures characteristic of different pneumonia subtypes. The four GLCM descriptors computed are:

$$\text{Contrast} = \sum_{i,j} |i - j|^2 P(i, j) \quad (2)$$

$$\text{Homogeneity} = \sum_{i,j} \frac{P(i, j)}{1 + |i - j|} \quad (3)$$

$$\text{Energy} = \sum_{i,j} P(i, j)^2 \quad (4)$$

$$\text{Correlation} = \sum_{i,j} \frac{(i - \mu_i)(j - \mu_j) P(i, j)}{\sigma_i \sigma_j} \quad (5)$$

Structural edge features are derived using Canny edge detection to trace pulmonary margins, extracting overall edge density and mean edge intensity as discriminative descriptors.

Global histogram features encode a 32-bin discrete intensity histogram, profiling the continuous global luminance distribution of each radiograph.

2.4 Dimensionality Reduction and Class Balancing

The 89-dimensional feature vector is subsequently compressed via Principal Component Analysis (PCA), retaining exactly 95% of cumulative explained variance, which empirically corresponds to 25–30 principal components. This step directly addresses the curse of dimensionality: initial experiments with raw pixel flattening yielded training accuracies approaching 99% alongside poor generalisation, validating the necessity of explicit dimensionality reduction.

Class imbalance is mitigated through the Synthetic Minority Over-sampling Technique (SMOTE) [5], applied exclusively to the training subset with k -neighbours set to 5. SMOTE synthesises realistic feature vectors for the minority (normal) class, restoring the training distribution to a balanced 1:1 ratio without discarding majority-class samples.

3 Literature Review

3.1 Handcrafted Feature Pipelines in Radiographic Analysis

Traditional computer-aided detection (CAD) systems follow a well-established sequential paradigm of preprocessing, explicit feature extraction, and classical classification. Within this paradigm, GLCM has historically served as a cornerstone for biomedical image classification, demonstrating consistent efficacy in distinguishing viral interstitial patterns from bacterial lobar consolidations [4]. Hybrid approaches that combine GLCM with Local Binary Patterns (LBP) and Histogram of Oriented Gradients (HOG) alongside Support Vector Machine classifiers have reported competitive accuracy on standardised radiographic benchmarks. The Canny edge detector has similarly been identified as optimal for radiological segmentation owing to its precision in delineating pulmonary opacity margins.

3.2 Managing High Dimensionality and Class Imbalance

PCA is widely established as the premier technique for condensing high-dimensional image feature vectors while preserving critical variance structure. SMOTE has become the industry standard for handling skewed radiographic datasets, with studies demonstrating that its integration with ensemble classifiers such as Random Forests stabilises learning dynamics, improves minority-class AUC, and elevates overall F1-scores [5].

3.3 Empirical Insights Informing the Design

The design decisions of the proposed architecture emerged directly from iterative experimentation. Applying classifiers to raw flattened pixel representations — tens of thousands of features — consistently produced severe overfitting, empirically confirming the necessity of PCA and handcrafted feature engineering. Downsampling to 64×64 pixels was identified as the resolution threshold that preserves sufficient macro-structure for Canny and GLCM analysis while enabling practical execution times for computationally intensive classifiers such as KNN and Random Forest.

4 Commercial Market Analysis and Industry Relevance

4.1 Market Dynamics and Key Players

The translation of AI-assisted diagnosis from academic research to deployable clinical utility represents a rapidly expanding commercial ecosystem. The AI diagnostics market is projected to sustain double-digit compound annual growth over the next decade. Established players including Qure.ai, Enlitic, and Zebra Medical have secured FDA clearance for comprehensive chest X-ray analysis systems that function as computer-assisted detection (CADe) tools, automating radiological triage by prioritising critical cases for specialist review and neutralising diagnostic bottlenecks in high-volume clinical environments.

4.2 Teleradiology and Scalable Deployment

A particularly impactful application domain is teleradiology, where AI-assisted platforms connect hospitals in geographically underserved regions with global networks of specialised radiologists, delivering immediate second-opinion overlays on local imaging studies. This model has proven especially effective in bridging the diagnostic capacity gap in low-to-middle income countries.

4.3 Barriers to Market Penetration

Despite this momentum, substantial barriers obstruct broad adoption. Comprehensive on-premises deep learning suites remain financially prohibitive for low-resource clinics, while stringent data protection legislation (HIPAA, GDPR) creates legal friction for cloud-based image processing. A persistent lack of Electronic Health Record (EHR) standardisation further complicates the rollout of unified diagnostic tools across heterogeneous institutional environments.

The classical machine learning approach proposed in this work offers a compelling alternative: an auditable, resource-efficient system deployable entirely on consumer-grade hardware, inherently compatible with local data privacy protocols, and immediately actionable in the resource-constrained clinical settings where the global burden of pneumonia is most acute.

References

- [1] World Health Organization (WHO). (2022). *Pneumonia Fact Sheet*.
- [2] Rajpurkar, P., et al. (2017). CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays with Deep Learning. *arXiv preprint arXiv:1711.05225*.
- [3] Kermany, D. S., et al. (2018). Identifying Medical Diagnoses and Treatable Diseases by Image-Based Deep Learning. *Cell*, 172(5), 1122–1131.
- [4] Haralick, R. M., Shanmugam, K., & Dinstein, I. (1973). Textural Features for Image Classification. *IEEE Transactions on Systems, Man, and Cybernetics*, 3(6), 610–621.
- [5] Chawla, N. V., et al. (2002). SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research*, 16, 321–357.